

# Problem set 4

2025-10-05

In the next problem set, we plan to explore the relationship between COVID-19 death rates and vaccination rates across US states by visually examining their correlation. This analysis will involve gathering COVID-19 related data from the CDC's API and then extensively processing it to merge the various datasets. Since the population sizes of states vary significantly, we will focus on comparing rates rather than absolute numbers. To facilitate this, we will also source population data from the US Census to accurately calculate these rates.

In this problem set we will learn how to extract and wrangle data from the data US Census and CDC APIs.

1. Get an API key from the US Census at [https://api.census.gov/data/key\\_signup.html](https://api.census.gov/data/key_signup.html). You can't share this public key. But your code has to run on a TFs computer. Assume the TF will have a file in their working directory named `census-key.R` with the following one line of code:

```
census_key <- "A_CENSUS_KEY_THAT_WORKS"
```

Write a first line of code for your problem set that defines `census_key` by running the code in the file `census-key.R`.

```
source("census_key.R")
```

2. The [US Census API User Guide](#) provides details on how to leverage this valuable resource. We are interested in vintage population estimates for years 2021 and 2022. From the documentation we find that the *endpoint* is:

```
url <- "https://api.census.gov/data/2021/pep/population"
```

Use the `httr2` package to construct the following GET request.

```
https://api.census.gov/data/2021/pep/population?get=POP_2020,POP_2021,NAME&for=state:*&key=Y
```

Create an object called `request` of class `httr2_request` with this URL as an endpoint. Hint: Print out `request` to check that the URL matches what we want.

```
library(httr2)

request <- paste("https://api.census.gov/data/2021/pep/population?get=POP_2020,POP_2021,NAME&for=state:*&key=Y")
request
```

```
[1] "https://api.census.gov/data/2021/pep/population?get=POP_2020,POP_2021,NAME&for=state:*&key=Y"
```

3. Make a request to the US Census API using the `request` object. Save the response to an object named `response`. Check the response status of your request and make sure it was successful. You can learn about *status codes* [here](#).

```
response <- request(request) |>
  req_perform()
resp_status(response)
```

```
[1] 200
```

4. Use a function from the **httr2** package to determine the content type of your response.

```
resp_content_type(response)
```

```
[1] "application/json"
```

5. Use just one line of code and one function to extract the data into a matrix. Hints: 1) Use the `resp_body_json` function. 2) The first row of the matrix will be the variable names and this OK as we will fix in the next exercise.

```
population <- response |> resp_body_json(simplifyVector = TRUE)

population
```

	[,1]	[,2]	[,3]	[,4]
[1,]	"POP_2020"	"POP_2021"	"NAME"	"state"
[2,]	"3962031"	"3986639"	"Oklahoma"	"40"
[3,]	"1961455"	"1963692"	"Nebraska"	"31"
[4,]	"1451911"	"1441553"	"Hawaii"	"15"
[5,]	"887099"	"895376"	"South Dakota"	"46"
[6,]	"6920119"	"6975218"	"Tennessee"	"47"
[7,]	"3114071"	"3143991"	"Nevada"	"32"
[8,]	"2117566"	"2115877"	"New Mexico"	"35"
[9,]	"3188669"	"3193079"	"Iowa"	"19"
[10,]	"2935880"	"2934582"	"Kansas"	"20"
[11,]	"690093"	"670050"	"District of Columbia"	"11"
[12,]	"29217653"	"29527941"	"Texas"	"48"
[13,]	"6154481"	"6168187"	"Missouri"	"29"
[14,]	"3012232"	"3025891"	"Arkansas"	"05"
[15,]	"10067664"	"10050811"	"Michigan"	"26"
[16,]	"1377848"	"1388992"	"New Hampshire"	"33"
[17,]	"10457177"	"10551162"	"North Carolina"	"37"
[18,]	"11790587"	"11780017"	"Ohio"	"39"
[19,]	"5130729"	"5190705"	"South Carolina"	"45"
[20,]	"577267"	"578803"	"Wyoming"	"56"
[21,]	"39499738"	"39237836"	"California"	"06"
[22,]	"778962"	"774948"	"North Dakota"	"38"
[23,]	"4651203"	"4624047"	"Louisiana"	"22"
[24,]	"6172679"	"6165129"	"Maryland"	"24"
[25,]	"991886"	"1003384"	"Delaware"	"10"
[26,]	"12989625"	"12964056"	"Pennsylvania"	"42"
[27,]	"10725800"	"10799566"	"Georgia"	"13"
[28,]	"4241544"	"4246155"	"Oregon"	"41"
[29,]	"5707165"	"5707390"	"Minnesota"	"27"
[30,]	"5784308"	"5812069"	"Colorado"	"08"
[31,]	"9279743"	"9267130"	"New Jersey"	"34"
[32,]	"4503958"	"4509394"	"Kentucky"	"21"
[33,]	"7718785"	"7738692"	"Washington"	"53"
[34,]	"1362280"	"1372247"	"Maine"	"23"
[35,]	"642495"	"645570"	"Vermont"	"50"
[36,]	"1847772"	"1900923"	"Idaho"	"16"
[37,]	"6785644"	"6805985"	"Indiana"	"18"
[38,]	"1086193"	"1104271"	"Montana"	"30"
[39,]	"20154933"	"19835913"	"New York"	"36"
[40,]	"3281538"	"3263584"	"Puerto Rico"	"72"
[41,]	"3600260"	"3605597"	"Connecticut"	"09"
[42,]	"21569932"	"21781128"	"Florida"	"12"

[43,]	"8632044"	"8642274"	"Virginia"	"51"
[44,]	"7022220"	"6984723"	"Massachusetts"	"25"
[45,]	"12785245"	"12671469"	"Illinois"	"17"
[46,]	"2956870"	"2949965"	"Mississippi"	"28"
[47,]	"7177986"	"7276316"	"Arizona"	"04"
[48,]	"3281684"	"3337975"	"Utah"	"49"
[49,]	"5892323"	"5895908"	"Wisconsin"	"55"
[50,]	"5024803"	"5039877"	"Alabama"	"01"
[51,]	"1789798"	"1782959"	"West Virginia"	"54"
[52,]	"1096229"	"1095610"	"Rhode Island"	"44"
[53,]	"732441"	"732673"	"Alaska"	"02"

- Examine the `population` matrix you just created. Notice that 1) it is not tidy, 2) the column types are not what we want, and 3) the first row is a header. Convert `population` to a tidy dataset. Remove the state ID column and change the name of the column with state names to `state_name`. Add a column with state abbreviations called `state`. Make sure you assign the abbreviations for DC and PR correctly. Hint: Use the **janitor** package to make the first row the header.

```
library(tidyverse)
library(janitor)
#population <- population |> ## Use janitor row to names function
# convert to tibble
# remove stat column
# rename state column to state_name
# use pivot_longer to tidy
# remove POP_ from year
# parse all relevant columns to numeric
# add state abbreviations using state.abb variable
# use case_when to add abbreviations for DC and PR

population <- population |>
  row_to_names(row_number=1) |>
  as.tibble(population) |>
  select(-state) |> rename(state_name = NAME) |>
  pivot_longer(cols = starts_with("POP_"), names_to = "year", values_to = "population") |>
  mutate(year = str_remove(year, "POP_"), population = as.numeric(population)) |>
  mutate(state = case_when(
    state_name == "District of Columbia" ~ "DC",
    state_name == "Puerto Rico" ~ "PR",
    TRUE ~ state.abb[match(state_name, state.name)]
  ))
```

```
print(population)
```

```
# A tibble: 104 x 4
  state_name year population state
  <chr>      <chr>      <dbl> <chr>
1 Oklahoma  2020      3962031 OK
2 Oklahoma  2021      3986639 OK
3 Nebraska  2020      1961455 NE
4 Nebraska  2021      1963692 NE
5 Hawaii    2020      1451911 HI
6 Hawaii    2021      1441553 HI
7 South Dakota 2020       887099 SD
8 South Dakota 2021       895376 SD
9 Tennessee  2020      6920119 TN
10 Tennessee  2021      6975218 TN
# i 94 more rows
```

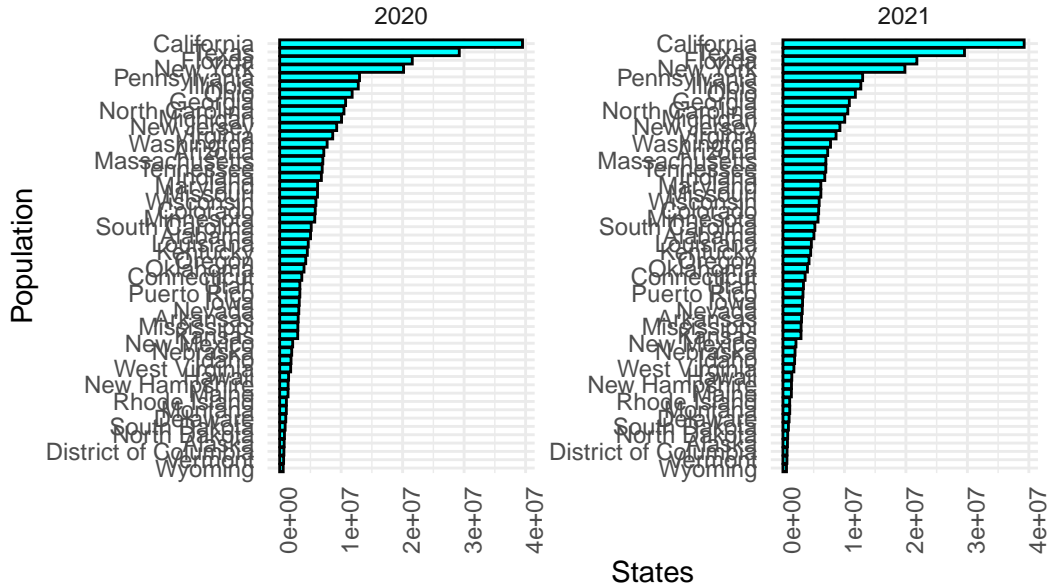
7. As a check, make a barplot of states' 2021 and 2022 populations. Show the state names in the y-axis ordered by population size. Hint: You will need to use `reorder` and use `facet_wrap`.

```
# population |>
# reorder state
# assign aesthetic mapping
# use geom_col to plot barplot
# flip coordinates
# facet by year

population.7 <- population |>
  mutate(state_name = reorder(state_name, population)) |>
  ggplot(aes(x = state_name, y = population)) +
  geom_col(fill = "cyan", color = "black") +
  coord_flip() +
  facet_wrap(~year, scales = "free_y") +
  labs(title = "State Populations for 2021 and 2022",
       x = "Population",
       y = "States") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))

population.7
```

## State Populations for 2021 and 2022



8. The following URL:

```
url <- "https://github.com/datasciencelabs/2025/raw/refs/heads/main/data/regions.json"
```

points to a JSON file that lists the states in the 10 Public Health Service (PHS) defined by CDC. We want to add these regions to the `population` dataset. To facilitate this create a data frame called `regions` that has two columns `state_name`, `region`, `region_name`. One of the regions has a long name. Change it to something shorter.

```
library(jsonlite)
library(purrr)
url <- "https://github.com/datasciencelabs/2025/raw/refs/heads/main/data/regions.json"
# regions <- use jsonlit JSON parser
# regions <- convert list to data frame. You can use map_df in purrr package

regions <- fromJSON(url, simplifyDataFrame = FALSE)

regions <- map_df(regions, function(x)
  data.frame(region=x$region, region_name = x$region_name, state_name = x$states))

regions$region_name <- gsub("New York and New Jersey, Puerto Rico, Virgin Islands", "NY & NJ", regions$region_name)
```

	region	region_name	state_name
1	1	New England	Connecticut
2	1	New England	Maine
3	1	New England	Massachusetts
4	1	New England	New Hampshire
5	1	New England	Rhode Island
6	1	New England	Vermont
7	2	NY & NJ & PR & VI	New Jersey
8	2	NY & NJ & PR & VI	New York
9	2	NY & NJ & PR & VI	Puerto Rico
10	2	NY & NJ & PR & VI	Virgin Islands
11	3	Mid-Atlantic	Delaware
12	3	Mid-Atlantic	District of Columbia
13	3	Mid-Atlantic	Maryland
14	3	Mid-Atlantic	Pennsylvania
15	3	Mid-Atlantic	Virginia
16	3	Mid-Atlantic	West Virginia
17	4	Southeast	Alabama
18	4	Southeast	Florida
19	4	Southeast	Georgia
20	4	Southeast	Kentucky
21	4	Southeast	Mississippi
22	4	Southeast	North Carolina
23	4	Southeast	South Carolina
24	4	Southeast	Tennessee
25	5	Midwest	Illinois
26	5	Midwest	Indiana
27	5	Midwest	Michigan
28	5	Midwest	Minnesota
29	5	Midwest	Ohio
30	5	Midwest	Wisconsin
31	6	South Central	Arkansas
32	6	South Central	Louisiana
33	6	South Central	New Mexico
34	6	South Central	Oklahoma
35	6	South Central	Texas
36	7	Central Plains	Iowa
37	7	Central Plains	Kansas
38	7	Central Plains	Missouri
39	7	Central Plains	Nebraska
40	8	Mountain States	Colorado
41	8	Mountain States	Montana
42	8	Mountain States	North Dakota

43	8	Mountain States	South Dakota
44	8	Mountain States	Utah
45	8	Mountain States	Wyoming
46	9	Pacific	Arizona
47	9	Pacific	California
48	9	Pacific	Hawaii
49	9	Pacific	Nevada
50	9	Pacific	American Samoa
51	9	Pacific	Commonwealth of the Northern Mariana Islands
52	9	Pacific	Federated States of Micronesia
53	9	Pacific	Guam
54	9	Pacific	Marshall Islands
55	9	Pacific	Republic of Palau
56	10	Pacific Northwest	Alaska
57	10	Pacific Northwest	Idaho
58	10	Pacific Northwest	Oregon
59	10	Pacific Northwest	Washington

9. Add a region and region name columns to the population data frame.

```
population <- population |>
  left_join(regions, by = "state_name")

population
```

# A tibble: 104 x 6

	state_name	year	population	state	region	region_name
	<chr>	<chr>	<dbl>	<chr>	<int>	<chr>
1	Oklahoma	2020	3962031	OK	6	South Central
2	Oklahoma	2021	3986639	OK	6	South Central
3	Nebraska	2020	1961455	NE	7	Central Plains
4	Nebraska	2021	1963692	NE	7	Central Plains
5	Hawaii	2020	1451911	HI	9	Pacific
6	Hawaii	2021	1441553	HI	9	Pacific
7	South Dakota	2020	887099	SD	8	Mountain States
8	South Dakota	2021	895376	SD	8	Mountain States
9	Tennessee	2020	6920119	TN	4	Southeast
10	Tennessee	2021	6975218	TN	4	Southeast

# i 94 more rows

10. From reading <https://data.cdc.gov/> we learn the endpoint <https://data.cdc.gov/resource/pwn4-m3yp.json> provides state level data from SARS-COV2 cases. Use the **httr2** tools you



have learned to download this into a data frame. Is all the data there? If not, comment on why.

```
api <- "https://data.cdc.gov/resource/pwn4-m3yp.json"
```

```
response <- request(api) |>  
  req_perform()
```

```
cases_raw <- response |>  
  resp_body_json(simplifyVector = TRUE) |>  
  as_tibble()
```

```
cases_raw
```

```
# A tibble: 1,000 x 10
```

	date_updated <chr>	state <chr>	start_date <chr>	end_date <chr>	tot_cases <chr>	new_cases <chr>	tot_deaths <chr>
1	2023-02-23T00:00:00~	AZ	2023-02-1~	2023-02~	2434631.0	3716.0	33042.0
2	2022-12-22T00:00:00~	LA	2022-12-1~	2022-12~	1507707.0	4041.0	18345.0
3	2023-02-23T00:00:00~	GA	2023-02-1~	2023-02~	3061141.0	5298.0	42324.0
4	2023-03-30T00:00:00~	LA	2023-03-2~	2023-03~	1588259.0	2203.0	18858.0
5	2023-02-02T00:00:00~	LA	2023-01-2~	2023-02~	1548508.0	5725.0	18572.0
6	2023-03-23T00:00:00~	LA	2023-03-1~	2023-03~	1580709.0	1961.0	18835.0
7	2023-04-27T00:00:00~	LA	2023-04-2~	2023-04~	1597070.0	1884.0	18937.0
8	2023-03-16T00:00:00~	NV	2023-03-0~	2023-03~	891702.0	1233.0	11937.0
9	2023-05-11T00:00:00~	FL	2023-05-0~	2023-05~	7572282.0	6937.0	88248.0
10	2022-10-27T00:00:00~	NYC	2022-10-2~	2022-10~	2928439.0	14590.0	42863.0

```
# i 990 more rows
```

```
# i 3 more variables: new_deaths <chr>, new_historic_cases <chr>,
```

```
#   new_historic_deaths <chr>
```

No, the data are not all there. Since the default limit is 1000 rows, we need to change the limit to a larger number.

We see exactly 1,000 rows. We should be seeing over  $52 \times 3$  rows per state.

11. The reason you see exactly 1,000 rows is because CDC has a default limit. You can change this limit by adding `$limit=10000000000` to the request. Rewrite the previous request to ensure that you receive all the data. Then wrangle the resulting data frame to produce a data frame with columns `state`, `date` (should be the end date) and `cases`. Make sure the cases are numeric and the dates are in `Date` ISO-8601 format.

```

api <- "https://data.cdc.gov/resource/pwn4-m3yp.json"

response <- request(api) |>
  req_url_query(`$limit` = 10000000000) |>
  req_perform()

cases_raw <- response |>
  resp_body_json(simplifyVector = TRUE) |>
  as_tibble()

#cases_raw

cases <- cases_raw |>
  select(state, end_date, new_cases) |>
  rename(date = end_date, cases = new_cases) |>
  mutate(
    cases = as.numeric(cases),
    date = as.Date(date))

cases

```

```

# A tibble: 10,380 x 3
   state date      cases
   <chr> <date>    <dbl>
1 AZ    2023-02-22  3716
2 LA    2022-12-21  4041
3 GA    2023-02-22  5298
4 LA    2023-03-29  2203
5 LA    2023-02-01  5725
6 LA    2023-03-22  1961
7 LA    2023-04-26  1884
8 NV    2023-03-15  1233
9 FL    2023-05-10  6937
10 NYC  2022-10-26 14590
# i 10,370 more rows

```

12. For 2020 and 2021, make a time series plot of cases per 100,000 versus time for each state. Stratify the plot by region name. Make sure to label you graph appropriately.

```

library(tidyverse)
library(lubridate)

```

```
cases <- cases |>
  filter(year(date) %in% c(2020, 2021)) |>
  left_join(population, by = "state") |>
  mutate(cases_per_100k = (cases / population) * 100000) |>
  filter(!is.na(cases_per_100k))
```

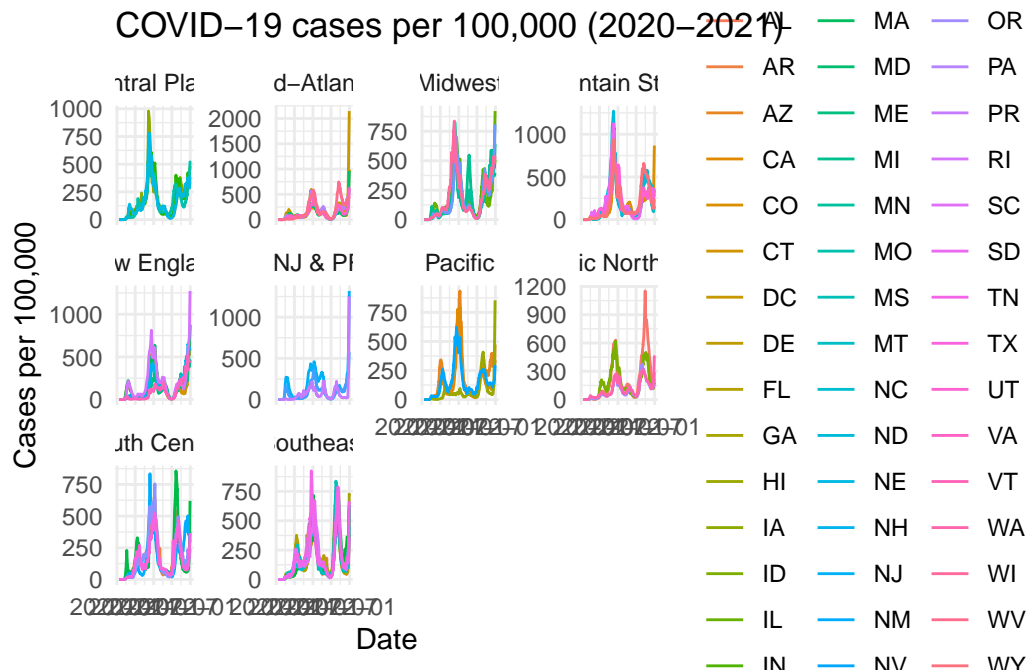
Warning in left\_join(filter(cases, year(date) %in% c(2020, 2021)), population, : Detected an to-many relationship between `x` and `y`.

i Row 1 of `x` matches multiple rows in `y`.

i Row 103 of `y` matches multiple rows in `x`.

i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to silence this warning.

```
cases |>
  ggplot(aes(x = date, y = cases_per_100k, color = state)) +
  geom_line() +
  facet_wrap(~ region_name, scales = "free_y") +
  labs(
    title = "COVID-19 cases per 100,000 (2020-2021)",
    x = "Date",
    y = "Cases per 100,000",
    color = "State"
  ) +
  theme_minimal()
```



- The dates in the `cases` dataset are stored as character strings. Use the **lubridate** package to properly parse the `date` column, then create a summary table showing the total COVID-19 cases by month and year for 2020 and 2021. The table should have columns for year, month (as month name), and total cases across all states. Order by year and month. Use the **knitr** package and `kable()` function to display the results as a formatted table.

```
library(lubridate)
library(knitr)

cases |>
  mutate(
    date = ymd(date),
    year = year(date),
    m_num = month(date),
    month = month(date, label = TRUE, abbr = FALSE)
  ) |>
  filter(year %in% c(2020, 2021)) |>
  group_by(year, month, m_num) |>
  summarise(total_cases = sum(cases, na.rm = TRUE), .groups = "drop") |>
  arrange(year, m_num) |>
  select(year, month, total_cases) |>
  kable()
```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.  
 Use 'xfun::attr2()' instead.  
 See help("Deprecated")

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.  
 Use 'xfun::attr2()' instead.  
 See help("Deprecated")

year	month	total_cases
2020	January	22
2020	February	136
2020	March	100670
2020	April	1645296
2020	May	1233382
2020	June	1285104
2020	July	3954032
2020	August	2904786
2020	September	2803834
2020	October	3217864
2020	November	7774444
2020	December	13815080
2021	January	11298230
2021	February	5087928
2021	March	3857498
2021	April	3388378
2021	May	1897906
2021	June	969634
2021	July	2241878
2021	August	7038814
2021	September	9921614
2021	October	4635708
2021	November	4578236
2021	December	10586782

14. The following URL provides additional COVID-19 data from the CDC in JSON format:

```
deaths_url <- "https://data.cdc.gov/resource/9bhg-hcku.json"
```

Use **httr2** to download COVID-19 death data from this endpoint. Make sure to remove the default limit to get all available data. Create a clean dataset called **deaths** with columns **state**,

date, and deaths (renamed from the original column name). Ensure dates are in proper Date format and deaths are numeric.

```
library(httr2)
library(tidyverse)
library(janitor)
library(lubridate)
api <- "https://data.cdc.gov/resource/9bhg-hcku.json"

response <- request(api) |>
  req_url_query(`$limit` = 1000000000) |>
  req_perform()

death_raw <- response |>
  resp_body_json(simplifyVector = TRUE) |>
  as_tibble()

#death_raw

deaths <- death_raw |>
  select(state, end_date, total_deaths) |>
  rename(date = end_date, deaths = total_deaths) |>
  mutate(
    deaths = as.numeric(deaths),
    date = as.Date(date)) |>
  filter(!state %in% c("United States", "Puerto Rico", "New York City", "District of Columbia"))

deaths
```

```
# A tibble: 127,500 x 3
```

	state	date	deaths
	<chr>	<date>	<dbl>
1	Alabama	2023-09-23	231602
2	Alabama	2023-09-23	1491
3	Alabama	2023-09-23	2691
4	Alabama	2023-09-23	344
5	Alabama	2023-09-23	453
6	Alabama	2023-09-23	2672
7	Alabama	2023-09-23	4549
8	Alabama	2023-09-23	5388
9	Alabama	2023-09-23	6827
10	Alabama	2023-09-23	8639

```
# i 127,490 more rows
```

15. Using the `deaths` dataset you created, make a bar plot showing the total COVID-19 deaths by state. Show only the top 10 states with the highest death counts. Order the bars from highest to lowest and use appropriate labels and title.

```
deaths_top10 <- deaths |>
  group_by(state) |>
  summarise(total_deaths = sum(deaths, na.rm = TRUE), .groups = "drop") |>
  arrange(desc(total_deaths)) |>
  slice_head(n = 10) |>
  mutate(state = factor(state, levels = state))

deaths_top10 |>
  ggplot(aes(x = state, y = total_deaths)) +
  geom_col(fill = "cyan", color = "black") +
  labs(
    title = "Top 10 Total COVID-19 Deaths by States",
    x = "States",
    y = "Total deaths"
  ) +
  theme_minimal()
```

